



**NTNU – Trondheim**  
Norwegian University of  
Science and Technology

## **OpenMP**

Guest lecture for TDT4200

Jan Christian Meyer 24.09.2024

# Pthreads reminder

- We start pthreads by giving them a function call to work on in the background
- When there's nothing left to do in the foreground, we wait for them to finish
- The threads have their own private values in the local variables of their function call
- The threads share the same copies of global variables, open files, *etc.*
  - We can coordinate access using *mutexes* (locks)

# A quick example

- I have prepared a short demonstration
  - It's in the code archive, under '01\_hello\_pthreads/'
- It starts 100 threads, each of them
  - starts with a brief, random wait (to upset the starting order)
  - takes a number from a global counter
  - increments the global counter
  - reports the number
- It uses a lock (mutex) to ensure that each thread gets a different number
  - If you remove the lock, there will be duplicates every now and again

# Things to do with pthreads

- Start
  - This requires a function call that takes a (void \*) argument, and returns a (void \*) result
- Stop
  - This happens when a thread's function call returns
- Synchronize using explicit locks
  - This requires the program to declare and manipulate each lock
- Send (and wait for) signals
  - If we attach a pthread\_cond\_t to a lock, threads that are waiting in line to lock it can be activated by the one that holds the lock
- Wait for everyone to arrive at a barrier
  - This is technically an optional pthreads feature, but it's often supported
  - When it's not, we can make barriers using locks+signals



# Pthreads code is verbose

- Every time we start a thread, it has to be
  - Declared
  - Initialized
  - Waited for
- Every time we want mutual exclusion, we need to
  - Declare a lock
  - Initialize it
  - Make sure that it's used correctly
  - Destroy it
- Every time we want a signal, we need to
  - Declare a condition variable
  - Initialize it
  - ...
- It's a lot of typing



# Pthreads are not (inherently) safe

- They just let you start several concurrent calls
  - If those calls write to the same locations in memory, they'll just do it in whatever order they are scheduled
  - This is not necessarily the same every time you run the program
- All synchronization is up to the programmer
  - Locks don't actually force threads to stay away from the section of code they protect
  - The threads must actively check the state of the lock and respect it in order to avoid race conditions
- This philosophy carries through to OpenMP
  - If you say something should be done in parallel when it should not, you will get wrong answers and no help from the compiler



# OpenMP

(Open Multi-Processing)

- OpenMP offers a less typing-intensive way to manage pthreads
  - OpenMP is not actually *required* to use pthreads behind the scenes, but practically every implementation does it
- It also offers a few other things, but thread control is a good place to start
  - ...and enough to cover in TDT4200



NTNU – Trondheim  
Norwegian University of  
Science and Technology

# What's in a name?

- We've heard about MPI
  - It's a specification for some function calls
  - Anyone can make their own implementation, so many exist
  - There's *mpich*, *IntelMPI*, *SpectrumMPI*, *mvapich*, ...
  - The idea is that you can switch variants without changing your code
- OpenMP is also a general specification with different implementations
  - There's *gomp*, *llvm/openmp*, *IOMP*, *xlc-openmp*, ...
  - Various flavors have come and gone since 1997
- In 2004, someone decided to call their MPI variant...  
(\*sigh\*) ... "OpenMPI"
  - I wish they would have chosen another name, but nobody asked me
  - It's easy to confuse the names, but please try not to, they are entirely different things
  - **OpenMP** is an interface for parallel programming with threads  
OpenMPI is one of many MPI alternatives, for parallel programming with processes





# OpenMP vs. pthreads

- If you want to, you can do the same things with OpenMP as you can with pthreads
- In today's code archive, '02\_hello\_openmp/' contains an OpenMP version of the same program as we used to demonstrate pthreads
- It's almost exactly the same, except that
  - The lock type is called `omp_lock_t`, and its init/destroy functions also have names with 'omp' in them
  - The lock/unlock functions are similarly renamed
  - The start/stop loops are gone – *what's this?*  
**`#pragma omp parallel num_threads(100)`**



# Before we proceed

- I just said that you can do the same things with OpenMP as with pthreads, but it's **not 100%** true:
  - You can't make threads go to sleep and then signal them to wake up again
  - The idea with OpenMP threads is that they will all have something to do for as long as they last, so there's no suspend/resume equivalent of the pthreads' condition variables
  - OpenMP threads are really easy to start and stop, so we can just stop unemployed threads, and start some new ones when there is extra work again

# Returning to our topic

- Instead of a 100-iteration loop that starts one thread per iteration, (and another that waits for them), we now have this:

```
#pragma omp parallel num_threads(100)  
hello_world();
```

- The same thing happens when the program runs, it's just shorter to write
- The compiler's preprocessor finds the `#pragma` directive, and inserts code to
  - spawn 100 threads
  - make each thread call 'hello\_world'
  - wait for them all to finish

...before the result is passed on to the compiler and treated as usual

# Changes in the Makefile

- There's another “small” difference:
  - The pthreads implementation has `-pthread` among the compiler flags
  - The OpenMP implementation has `-fopenmp` in the same place
- The difference is actually quite profound
  - The first one tells the compiler to find the pthreads library, and link it up with all the pthread function calls that we put in the source code
  - The second one tells the preprocessor to find these “`#pragma omp`” things in the source code, and use them to edit in a bunch of pthreads code that we didn't explicitly write



# So, what is this #pragma?

- Generally speaking,  
    #pragma is a way for compiler authors to extend the programming language  
    with features that may or may not be supported by other compilers
- If you put #pragma directives in your code that the  
    compiler doesn't understand, it will just ignore them
- They can request anything, you can literally write  
    #pragma make me a sandwich  
    in your program, and the code will work just fine  
    ...but you will only get a sandwich if you have a compiler that knows how to  
    make one
- Try it at home!



# The OpenMP solution

- Directives that begin with **#pragma omp** suggest that you hope your compiler is prepared to create some pthread code for you
- Pretty much all of the compilers used in HPC know how to do this, so it's OK to expect some automatic parallelization to happen
- If you wrote your code without *requiring* it to be parallelized, it will still work with compilers that don't understand OpenMP directives
  - It will just run a lot slower



# The anatomy of a directive

- All the OpenMP directives go

`#pragma omp <directive> [clauses]`

This says we're  
writing OpenMP

This says what  
kind of work we  
want done to  
the code

These are optional  
parts that adjust  
how the directive acts

- In the example, 'parallel' is the directive
  - It says that we want to start a bunch of threads
- `num_threads(#)` is a clause
  - If you don't specify it, you'll get one thread per core in the machine by default



# Private and shared variables

- We mentioned that pthreads have private variables inside the scope of the function they call

```
void *threaded_function ( void *args ) {  
    int my_value = rand();    // <--- only this thread has this value  
}
```

- This is because function calls allocate local variables in *stack memory*, and each thread has its own stack
- Function calls are actually a kind of special case for this, stack allocation happens more often than one might think



# Basic blocks

- When you put some statements between { curly braces } in C, you get an opportunity to declare variables that only last as long as the (basic) block
- A basic block is a statement, so you can put it almost anywhere
- Witness:

```
int main() {  
    int a = 64, b = 32, c;  
    {  
        int a = 42;                // override the a we had (i.e. push a different a on  
        top of stack)  
        c = a - b; // use the closest definition of a (i.e. the topmost one on the stack)  
    }  
    printf ( "%d\n", c );          // <--- this prints 10, even though a is 64 again now  
}
```



# Basic blocks with names

- Functions in C are pretty much just top-level basic blocks that have been decorated with a name
  - They also feature some local variables that can get values from outside (arguments)
  - They let you extract one local value and keep it for later (return value)
  - Otherwise, they're not much more than a bit of stack memory where we keep block-local declarations
- OpenMP uses this fact to let you spawn threads that last as long as a basic block
  - There's an example in the code archive, under `03_hello_basic_block`



# When there isn't a basic block

- In the second example, we have

```
#pragma omp parallel  
hello_world();
```

- This is (practically) equivalent to

```
#pragma omp parallel  
{  
    hello_world();  
}
```

- OpenMP directives apply to the next statement  
...but as we've seen, the next statement can be a block



# Identity numbers and such

- Inside a parallel region, each thread can obtain its index using  
`omp_get_thread_num()`  
and find out how many active threads there are with  
`omp_get_num_threads()`
- Outside of parallel regions, these will respond with 0 and 1
  - If you want to know how many threads will be launched without actually starting them, there is  
`omp_get_max_threads()`



# Now we can do All The Things™

- As with MPI, all we really need is
  - An ID number
  - Size of the collective
  - A way to communicate (threads can use shared data + locks)...and then all kinds of parallel things can happen
- The rest of OpenMP is window dressing to make stuff easier to write
  - It's really practical, though - hardly anyone needs to do the thread-id and locking thing explicitly, I just wanted to show that it's there
  - Next lecture, we'll talk about the *worksharing directives*