



NTNU – Trondheim
Norwegian University of
Science and Technology

OpenMP worksharing directives

Guest lecture for TDT4200

Jan Christian Meyer 24.09.2024

Partitioning shared work

- We have seen that we can assign work to threads based on their index in the thread pool:

```
int tid = omp_get_thread_num();
```

- This is a little bit of a hassle

- For thread-specific blocks of code, we need something like this

```
if ( tid == 0 ) { /* Do one thing */ }  
else if ( tid == 1 ) { /* Do another thing */ }
```

...

- For loops, we need to combine the index with the induction variable to work out a selection of iterations

```
for ( int x=tid; x<x_max; x+=n_threads )    ← round robin  
for ( int x=bottom[tid]; x<top[tid]; x++ )  ← consecutive range
```

- It is not super difficult, but it's repetitive to type every time
 - Also extremely common, so it can be automated



Worksharing directives to the rescue!

- These are OpenMP directives that can split a given workload between threads for you, without requiring you to do anything based on the thread id#
- We'll look at three flavors
 - Sections
 - Loops
 - Single



Functional decomposition

- This is when we split the work by the function of its sub-tasks
 - We've talked about it in terms of *pipelining*

This thing only
installs seats →



← This thing only
slaps on doors

Throughput doubles
when the pipeline is full

Partial products roll past in this direction



NTNU – Trondheim
Norwegian University of
Science and Technology

Data decomposition

- This is when we split the work by the input/output of its sub-tasks
 - Pretty much all we've been doing so far, because you don't have to design additional code in order to increase the number of participants

Everyone does
the same thing →



↑ ↑ ↑ ↑ ↑ ↑ ↑ ↑
They do it to individually assigned parts of a field



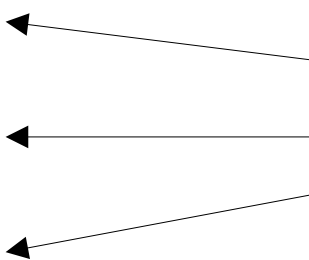
NTNU – Trondheim
Norwegian University of
Science and Technology

Sections

- For functional decomposition, OpenMP has *sections*

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        { /* Section #1 */ }
        #pragma omp section
        { /* Section #2 */ }
        #pragma omp section
        { /* Section #3 */ }
    }
}
```

Each of these
will be run by
one thread



Implicit synchronization

- Worksharing directives have an implicit barrier at the end

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        { /* Section #1 */ }
        #pragma omp section
        { /* Section #2 */ }
        #pragma omp section
        { /* Section #3 */ }
    }
}
```

All threads will
synchronize here
by default



Implicit synchronization

- Because they very often occur just after each other

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        { /* Section #1 */ }
        #pragma omp section
        { /* Section #2 */ }
    }
    #pragma omp sections
    {
        #pragma omp section
        { /* Section #3 */ }
        #pragma omp section
        { /* Section #4 */ }
    }
}
```

Finish these first

Synchronize

Finish these next

Synchronize again

Implicit barriers make these two blocks of sections work as separate stages



Clauses

- Most OpenMP directives have an optional set of additional terms that can control details of their semantics
 - We've already seen the *num_threads* clause for the *parallel* directive
- The worksharing directives have a clause *nowait*
 - Its use indicates that you wish to omit the implicit barrier at the end



nowait in practice

- The nr. of sections limits the number of threads in use, additional threads wait

```
#pragma omp parallel  
{
```

```
  #pragma omp sections  
  {
```

```
    #pragma omp section
```

```
    { /* Section #1 */ }
```

```
    #pragma omp section
```

```
    { /* Section #2 */ }
```

```
  }
```

```
  #pragma omp sections
```

```
  {
```

```
    #pragma omp section
```

```
    { /* Section #3 */ }
```

```
    #pragma omp section
```

```
    { /* Section #4 */ }
```

```
  }
```

```
}
```

Two threads here

Stop

Two threads here

By default, this example
will only use 2 threads
at a time



nowait in practice

- If we omit the implicit barrier, additional threads will “fall through” and start working

```
#pragma omp parallel
```

```
{
```

```
  #pragma omp sections nowait
```

```
  {
```

```
    #pragma omp section
```

```
    { /* Section #1 */ }
```

```
    #pragma omp section
```

```
    { /* Section #2 */ }
```

```
  }
```

```
  #pragma omp sections
```

```
  {
```

```
    #pragma omp section
```

```
    { /* Section #3 */ }
```

```
    #pragma omp section
```

```
    { /* Section #4 */ }
```

```
  }
```

```
}
```

Skip the barrier

Two threads here

Two more threads here, right away

Here, we have enough sections to employ 4 threads at a time



That's a silly example

- Yes, it is.
 - A simpler way to write the same effect would be to just include all four sections under the same `#pragma omp sections` directive to begin with
- I just wanted to make a simple illustration of the `nowait` clause
 - It applies to the other worksharing directives as well
 - It's occasionally useful



Saving some keystrokes

- A very common use case is to start some threads only to give them exactly 1 worksharing directive

- With *sections* as an example, it creates this pattern

```
#pragma omp parallel
{
    #pragma omp sections
    {
        ...
    }
}
```

- Because it's redundant to separate the thread starting/stopping directive from the work partitioning when there's only 1, we can write them together

```
#pragma omp parallel sections
{
    ...
}
```

- This means exactly the same thing as above

(But there will always be an implicit synch. at the end, because the threads join there)



Loops

- Parallelizing loops by thread indices amounts to partitioning its *iteration space*

for (i=tid; i<N; i+=n_threads)

- assigns every $(n_threads)^{th}$ iteration to a thread

for (i=bottom[tid]; i<top[tid]; i++)

- assigns blocks of $top[tid]-bottom[tid]$ iterations to a thread

- When we have a loop with an induction variable (such as for loops in C) this assignment can be done automatically

```
#pragma omp parallel for
```

```
for ( int i=0; i<N; i++ )
```

makes some default mapping of iterations to threads

(Note that we didn't *have* to join the “parallel” and “for” parts, you can also have several instances of `#pragma omp for` inside one `#pragma omp parallel`)

- There's no equivalent for *while* loops, because we can't predict their iteration counts in the same way

(There's another technique, but we'll get back to it)



The *schedule*(kind,blocksize) clause

- Parallel loop directives allow you to control how they partition the iterations between threads
- The *blocksize* is an optional, positive integer which we shall discuss imminently
- The *kind* is one of these:
 - static
 - dynamic
 - guided
 - auto
 - runtime



A unit of work

- When OpenMP partitions an iteration space between threads, it has some leeway with how many iterations to include in “one work unit”
- The absolutely smallest unit available/possible is to distribute 1 iteration at a time
 - Units of 1 iteration gives the round-robin assignment we’ve worked out manually
- Depending on how much work each iteration contains, 1 iteration can easily be a bit on the short side
- Increasing the unit size makes the work distribution more coarse-grained



Small vs. big work units

- Big blocks:
 - Fewer units to distribute, and hence, less scheduling to do
BUT
 - There's a limit to how big the blocks should be
 - At the extreme end: if the entire iteration space is one big block, we've taken away all the parallelism again
- Small blocks:
 - More units to distribute, more disruptions in memory access pattern
BUT
 - Greater flexibility to assign work to unemployed threads

The block size

- It's easy to assume that the block size parameter of the schedule is the number of iterations handed to each thread
- This is not always true
- It is the minimal number of iterations handed to a thread
 - Some of the schedule kinds take the liberty to hand out bigger blocks
 - They won't hand out smaller blocks if they can help it, though
 - It's intended to be a measure of how few iterations it can make sense to lump together
 - This number depends on the details of your program, so you can set it



Automatic schedule

- This is the default kind
- It doesn't have to be a particular fixed kind, it's the one that your OpenMP implementation nominates as most likely to do the best job in the greatest number of cases
- It tends to be a good guess for nested loops that sweep over multidimensional arrays with approximately equal workloads per element
 - That's a very common use case



Runtime schedule

- This is for when you don't want to embed the choice of schedule into your program
- With a runtime schedule, your OpenMP program will search the calling shell's environment variables for a specification of what to use on each run:

```
export OMP_SCHEDULE="static,4"
```

```
./my_program    # program runs with schedule(static,4) as default
```

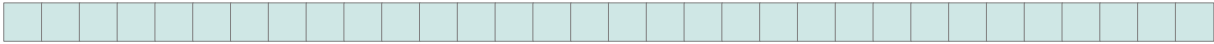
```
export OMP_SCHEDULE="dynamic,16"
```

```
./my_program    # program runs with schedule(dynamic,16) instead
```



Static schedule

- This is the schedule we've been calculating by hand throughout all this threading stuff
- Given an iteration space

0  max

schedule(static,6) will assign iterations to e.g. threads
0, 1, and 2, thus:

0  max



The master/worker pattern

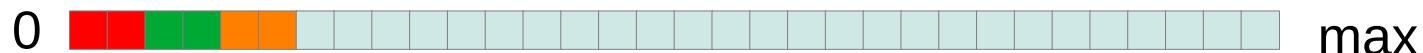
- We haven't made much use of it, but a common way to implement work sharing in queued systems is to
 - Keep an active *thread pool* of available worker threads
 - Keep a queue of finite work packages (which may or may not grow/shrink while the program is running)
 - Assign the next package in the queue every time a worker thread becomes available
- Web servers, transactional databases, and other on-line request processing systems tend to do this
 - HPC programs rarely have infinite streams of incoming requests
 - They still use this pattern to achieve some measure of *load balancing* when the amount of work in each package is unevenly distributed



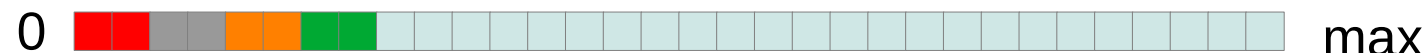
Dynamic schedule

- The dynamic kind of schedule works this way
- We can illustrate `schedule(dynamic,2)` with our iteration space and 3 threads again:

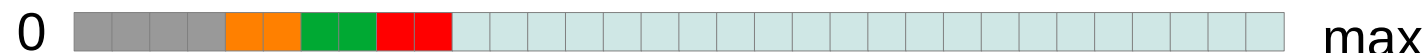
Everyone gets something to begin with



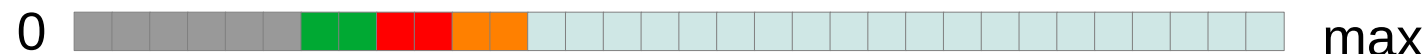
If **thread 1** finishes first, it gets the next unit



Thread 0 may be the next one out



By now, **thread 2** may have come around



...and so, it continues...



NTNU – Trondheim
Norwegian University of
Science and Technology

Guided schedule

- This one is similar to dynamic, but acknowledges the observation that we probably have a barrier coming up at the end of the loop
- While the barrier remains far in the future, it doesn't matter so much how big the blocks are
 - Workers that run out of work can just pick up some more
- When the barrier is imminent, it's a mistake to hand out a giant workload to one worker
 - Everyone else will have to wait for it to finish
- Guided schedule starts with big blocks and gradually shrink them down to the blocksize

schedule(guided,1) illustrated

- Everyone gets lots of work at the beginning:



- Past the halfway mark, we should probably shrink the workloads we dispense:

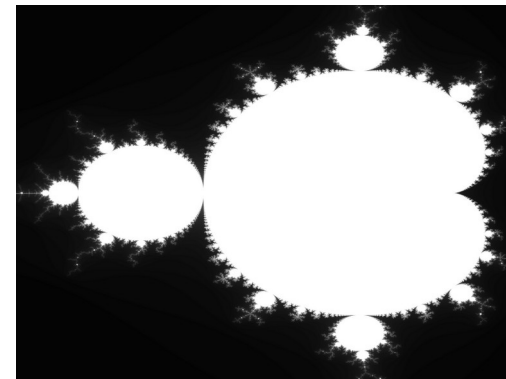


- Near the end, everyone gets the smallest available block sizes, to minimize the inevitable wait



Demonstration time

- Today's code archive has a subdirectory 01_fractal
 - It has an OpenMP-enabled Mandelbrot set generator in it
- We don't have to dwell on the calculation, but the output image looks like this:
 - The important characteristic is that a black point represents a loop that has terminated immediately, while a white point has required 256 iterations
- The loop over the y-axis is the parallel one
- Clearly, some horizontal lines contain much more work than others



The experiment

- The program times its own execution
 - Conveniently, this also lets us demonstrate the function `double omp_get_wtime (void);`
 - It's exactly like `MPI_Wtime()`, in that it returns some number of seconds
 - It's also exactly like `MPI_Wtime()` in that implementations tend to use precisely the same system clock
- If our theory is correct, this program might run faster with a guided schedule than with the automatic
 - Let's try it out

Single (and master)

- If you are inside a parallel region but only want one thread to do something, you can mark a block with
`#pragma omp single`
and only one thread (but *any* thread) will go in there
- There's another flavor called
`#pragma omp master`
which makes sure that only thread 0 goes in there
- It's a flavor of mutual exclusion without having to declare locks

